Bachelor's Thesis

Financial prediction with Neural Networks "Predicting Spread Prices on VIX Futures with Deep Learning"

by Thomas Leyh

University of Freiburg Department of Computer Science Computer Vision Group Prof. Dr. Thomas Brox

September 12th, 2017

Outline

- Why do this?
 - Some financial terminology
 - The actual problem
 - Where machine learning comes in
- 2 Approaches
 - Data representation
 - Network architecture
- Conclusion
 - Results
 - Why did it fail?
 - Where to go from here?

Outline

- Why do this?
 - Some financial terminology
 - The actual problem
 - Where machine learning comes in
- 2 Approaches
 - Data representation
 - Network architecture
- 3 Conclusion
 - Results
 - Why did it fail?
 - Where to go from here?

the S&P 500 Index (similar to the German DAX).

Volatility Measures the risk of sharp market movements.

Futures A standardized tradable contract
The holder is guaranteed the
delivery of the underlying
commodity at its expiration.

VIX CBOE's Volatility Index based on the S&P 500 Index (similar to the German DAX).

Volatility Measures the risk of sharp market movements.

Futures A standardized tradable contract
The holder is guaranteed the
delivery of the underlying
commodity at its expiration.

VIX CBOE's Volatility Index based on the S&P 500 Index (similar to the German DAX).

Volatility Measures the risk of sharp market movements.

Futures A standardized tradable contract.

The holder is guaranteed the delivery of the underlying commodity at its expiration.

VIX CBOE's Volatility Index based on the S&P 500 Index (similar to the German DAX).

Volatility Measures the risk of sharp market movements.

Futures A standardized tradable contract.

The holder is guaranteed the delivery of the underlying commodity at its expiration.



Source: Jean Weber, INRA (Wikimedia Commons)

VIX CBOE's Volatility Index based on the S&P 500 Index (similar to the German DAX).

Volatility Measures the risk of sharp market movements.

Futures A standardized tradable contract.

The holder is guaranteed the delivery of the underlying commodity at its expiration.

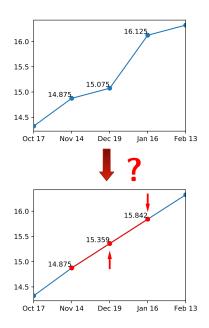


Source: Jean Weber, INRA (Wikimedia Commons)

The actual problem

We are interested in the correction of the term structure:

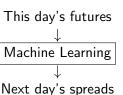
- Will it get corrected?Do we invest in the first place?
- At what position? Where to place the spread?
- When will the correction occur? How long to hold the spread?



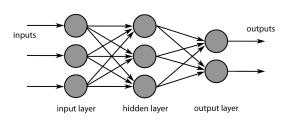
Where machine learning comes in

We want some working model for this prediction. But we are too lazy to build it ourselves.

- \rightarrow Let the computer do it.
 - Select appropriate model family
 - Optimize hyperparameters
 - Throw data at it
 - 4 . . .
 - Profit



Where machine learning comes in: Neural Networks



Source: Chrislb (Wikimedia Commons)

Equivalent to:

$$f^{(1)} \begin{pmatrix} \begin{bmatrix} y_1^{(0)} \\ y_2^{(0)} \\ y_3^{(0)} \end{bmatrix}^\top \cdot \begin{bmatrix} w_{1,1}^{(1)} & w_{1,2}^{(1)} & w_{1,3}^{(1)} \\ w_{2,1}^{(1)} & w_{2,2}^{(1)} & w_{2,3}^{(1)} \\ w_{3,1}^{(1)} & w_{3,2}^{(1)} & w_{3,3}^{(1)} \end{bmatrix} = y^{(1)\top}$$

$$f^{(2)} \left(y^{(1)\top} W^{(2)} \right) = y^{(2)\top}$$

$$= y^{(2)\top}$$

with often very simple activation function f working elementwise.

outputs

Outline

- Why do this?
 - Some financial terminology
 - The actual problem
 - Where machine learning comes in
- 2 Approaches
 - Data representation
 - Network architecture
- Conclusion
 - Results
 - Why did it fail?
 - Where to go from here?

Data representation

- There were 2656 samples used
 - From October 23th, 2006 to May 11th, 2017
 - ▶ Split with ratio 70:15:15 for training, validation and testing
 - Optionally using minutely data
- Preprocessing of the input data today's futures
 - ▶ The difference between the term structures' legs was used
 - ► A combination of additional inputs was used, mostly *days until expiration*
 - Optionally centered and normalized
- Choosing the target data tomorrow's spreads
 - Predict all spreads at once or
 - Predict one spread at a time or
 - Just predict if prices will fall, rise or stay the same

Data representation

- There were 2656 samples used
 - From October 23th, 2006 to May 11th, 2017
 - Split with ratio 70:15:15 for training, validation and testing
 - Optionally using minutely data
- Preprocessing of the input data today's futures
 - ▶ The difference between the term structures' legs was used
 - A combination of additional inputs was used, mostly days until expiration
 - Optionally centered and normalized
- Choosing the target data tomorrow's spreads
 - Predict all spreads at once or
 - Predict one spread at a time or
 - Just predict if prices will fall, rise or stay the same

Data representation

- There were 2656 samples used
 - From October 23th, 2006 to May 11th, 2017
 - ▶ Split with ratio 70:15:15 for training, validation and testing
 - Optionally using minutely data
- Preprocessing of the input data today's futures
 - ▶ The difference between the term structures' legs was used
 - A combination of additional inputs was used, mostly days until expiration
 - Optionally centered and normalized
- Choosing the target data tomorrow's spreads
 - Predict all spreads at once or
 - Predict one spread at a time or
 - Just predict if prices will fall, rise or stay the same

Network architecture

Using Feedforward Neural Networks.

A very basic model where layers are simply stacked on top of each other.

Hyperparameters:

Depth Number of hidden layers

Width Number of nodes per layer

Activations For introducing nonlinearity at each at each node's ouput

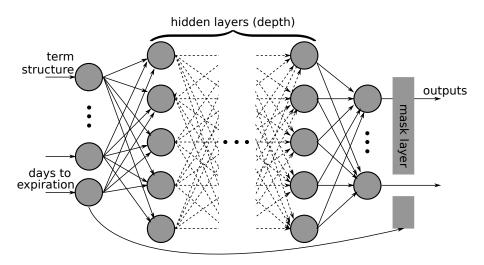
Loss function For comparing predictions and target values at the network's end (here: *mean squared error*)

Optimizer For iteratively updating the weights (mostly Adam)

Regularizations So the network doesn't just learn data by heart

... and many more ...

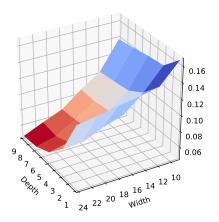
Network architecture



Outline

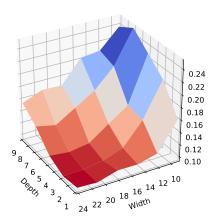
- Why do this?
 - Some financial terminology
 - The actual problem
 - Where machine learning comes in
- 2 Approaches
 - Data representation
 - Network architecture
- Conclusion
 - Results
 - Why did it fail?
 - Where to go from here?

During training the loss mostly looks like this:



But that's just overfitting – learning the training data by heart.

The loss when looking at validation data (unseen by network):



Deep networks have terrible performance on the problem.

After finding good hyperparameters, how do we do on the test set?

	Epoch	Test Loss
Naive		0.1318
Basic	307	0.1563
With dropout	174	0.2984
Self-normalizing	242	0.1259
Minutely data	70	0.1294
Additional inputs	268	0.1321
Single spreads	200	0.1760

After finding good hyperparameters, how do we do on the test set?

	Epoch	Test Loss
Naive	!!!	0.1318
Basic	307	0.1563
With dropout	174	0.2984
Self-normalizing	242	0.1259
Minutely data	70	0.1294
Additional inputs	268	0.1321
Single spreads	200	0.1760

Our best model has equal performance as the naive approach.



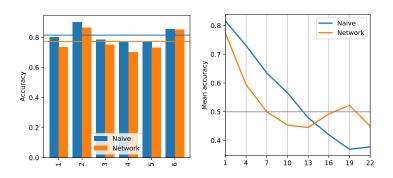


Let's make the problem easier: Just predict if spread prices will:

- (a) rise
- (b) fall
- (c) roughly stay the same

Let's make the problem easier: Just predict if spread prices will:

- (a) rise
- (b) fall
- (c) roughly stay the same



Conclusions:

- Neural networks work best in vision, speech and language
 - ► These are easy for humans
 - ▶ Very high-dimensional (many inputs)
- Financial problems are mostly hard for humans
- Additionally few inputs were used here
- Therefore the network only learned an (bad) approximation of the naive solution

"Most tasks that consist of mapping an input vector to an output vector, and that are easy for a person to do rapidly, can be accomplished via deep learning..."

Conclusions:

- Neural networks work best in vision, speech and language
 - ► These are easy for humans
 - Very high-dimensional (many inputs)
- Financial problems are mostly hard for humans
- Additionally few inputs were used here
- Therefore the network only learned an (bad) approximation of the naive solution

"Most tasks that consist of mapping an input vector to an output vector, and that are easy for a person to do rapidly, can be accomplished via deep learning..."

Conclusions:

- Neural networks work best in vision, speech and language
 - ► These are easy for humans
 - Very high-dimensional (many inputs)
- Financial problems are mostly hard for humans
- Additionally few inputs were used here
- Therefore the network only learned an (bad) approximation of the naive solution

"Most tasks that consist of mapping an input vector to an output vector, and that are easy for a person to do rapidly, can be accomplished via deep learning..."

Conclusions:

- Neural networks work best in vision, speech and language
 - ► These are easy for humans
 - Very high-dimensional (many inputs)
- Financial problems are mostly hard for humans
- Additionally few inputs were used here
- Therefore the network only learned an (bad) approximation of the naive solution

"Most tasks that consist of mapping an input vector to an output vector, and that are easy for a person to do rapidly, can be accomplished via deep learning..."

Conclusions:

- Neural networks work best in vision, speech and language
 - ▶ These are easy for humans
 - Very high-dimensional (many inputs)
- Financial problems are mostly hard for humans
- Additionally few inputs were used here
- Therefore the network only learned an (bad) approximation of the naive solution

"Most tasks that consist of mapping an input vector to an output vector, and that are easy for a person to do rapidly, can be accomplished via deep learning..."

Conclusions:

- Neural networks work best in vision, speech and language
 - ► These are easy for humans
 - Very high-dimensional (many inputs)
- Financial problems are mostly hard for humans
- Additionally few inputs were used here
- Therefore the network only learned an (bad) approximation of the naive solution

"Most tasks that consist of mapping an input vector to an output vector, and that are easy for a person to do rapidly, can be accomplished via deep learning..."

Conclusions:

- Neural networks work best in vision, speech and language
 - ► These are easy for humans
 - Very high-dimensional (many inputs)
- Financial problems are mostly hard for humans
- Additionally few inputs were used here
- Therefore the network only learned an (bad) approximation of the naive solution

"Most tasks that consist of mapping an input vector to an output vector, and that are easy for a person to do rapidly, can be accomplished via deep learning..."

Conclusions:

- Neural networks work best in vision, speech and language
 - ► These are easy for humans
 - Very high-dimensional (many inputs)
- Financial problems are mostly hard for humans
- Additionally few inputs were used here
- Therefore the network only learned an (bad) approximation of the naive solution

"Most tasks that consist of mapping an input vector to an output vector, and that are easy for a person to do rapidly, can be accomplished via deep learning..."

Where to go from here?

- Use other machine learning techniques when appropriate
- Neural networks for high-dimensional input data
 - Self-normalizing networks seem promising
- And most important: Learn more math

For code, data and further analysis see:

https://github.com/leyhline/vix-term-structure

Where to go from here?

- Use other machine learning techniques when appropriate
- Neural networks for high-dimensional input data
 - Self-normalizing networks seem promising
- And most important: Learn more math

For code, data and further analysis see:

https://github.com/leyhline/vix-term-structure

Thanks for listening.